

1. Import libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import (confusion_matrix, classification_report, accuracy_score,
                             precision_score, recall_score, f1_score)
```

2. Read Dataset (simulate data for this example) np.random.seed(42)

```
# 3. Read Dataset (simulate data for this example)
np.random.seed(42)
df = pd.DataFrame({
    'Age': np.random.randint(20, 80, 1000),
    'BMI': np.random.uniform(18, 35, 1000),
    'BloodPressure': np.random.randint(90, 180, 1000),
    'Glucose': np.random.randint(70, 200, 1000),
    'Gender': np.random.choice(['Male', 'Female'], 1000),
    'Geography': np.random.choice(['Urban', 'Rural'], 1000),
    'Smoker': np.random.choice([0, 1], 1000),
    'FamilyHistory': np.random.choice([0, 1], 1000),
    'DiseasePresent': np.random.choice([0, 1], 1000, p=[0.7, 0.3])
})
```

3. exploratory data analysis

```
# 4.1 Shape
print("Dataset Shape:", df.shape)
```

```
# 4.2 Preview
print(df.head())
```

```
# 4.3 Summary
print(df.info())
```

```
# 4.4 Statistical Properties
print(df.describe())
```

```
Dataset Shape: (1000, 9)
```

	Age	BMI	BloodPressure	Glucose	Gender	Geography	Smoker	\
0	58	18.797240	173	145	Female	Urban	1	
1	71	22.567432	157	145	Male	Rural	1	
2	48	18.377141	96	116	Female	Urban	1	
3	34	26.468808	143	84	Female	Rural	1	
4	62	26.095582	90	174	Female	Rural	1	

	FamilyHistory	DiseasePresent
0	0	0
1	0	0
2	1	0
3	1	0
4	1	0

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              1000 non-null   int64
1   BMI              1000 non-null   float64
2   BloodPressure    1000 non-null   int64
3   Glucose          1000 non-null   int64
4   Gender           1000 non-null   object
5   Geography        1000 non-null   object
6   Smoker           1000 non-null   int64
7   FamilyHistory    1000 non-null   int64
```

```

      8  DiseasePresent  1000 non-null   int64
dtypes: float64(1), int64(6), object(2)
memory usage: 70.4+ KB
None

```

	Age	BMI	BloodPressure	Glucose	Smoker \
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	50.200000	26.709609	132.908000	135.280000	0.484000
std	17.372905	4.827534	26.305542	36.699434	0.499994
min	20.000000	18.004038	90.000000	70.000000	0.000000
25%	36.000000	22.637515	110.000000	104.750000	0.000000
50%	51.000000	27.015266	131.000000	135.000000	0.000000
75%	66.000000	30.790958	156.000000	168.000000	1.000000
max	79.000000	34.989009	179.000000	199.000000	1.000000

	FamilyHistory	DiseasePresent
count	1000.000000	1000.000000
mean	0.528000	0.290000
std	0.499465	0.453989
min	0.000000	0.000000
25%	0.000000	0.000000
50%	1.000000	0.000000
75%	1.000000	1.000000
max	1.000000	1.000000

4. Feature Selection

```

features = df.drop('DiseasePresent', axis=1)
target = df['DiseasePresent']

```

5. Convert Categorical Columns to Numeric

```

# 6.1 Gender
df['Gender'] = LabelEncoder().fit_transform(df['Gender'])

# 6.2 Geography
df['Geography'] = LabelEncoder().fit_transform(df['Geography'])

```

6. Feature Scaling

```

X = df.drop('DiseasePresent', axis=1)
y = df['DiseasePresent']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

```


7. Model Training

```
# 8.1 Predict accuracy with different algorithms
models = {
    "Logistic Regression": LogisticRegression(),
    "Random Forest": RandomForestClassifier(),
    "Decision Tree": DecisionTreeClassifier(),
    "K-Nearest Neighbors": KNeighborsClassifier(),
    "Support Vector Machine": SVC()
}

accuracy_scores = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    accuracy_scores[name] = acc
    print(f"{name} Accuracy: {acc:.2f}")

# 8.2 Plot classifier accuracy scores
plt.figure(figsize=(10, 5))
sns.barplot(x=list(accuracy_scores.keys()), y=list(accuracy_scores.values()))
plt.xticks(rotation=45)
plt.title("Model Accuracy Comparison")
plt.ylabel("Accuracy")
plt.show()
```

```
Logistic Regression Accuracy: 0.73
Random Forest Accuracy: 0.69
Decision Tree Accuracy: 0.61
K-Nearest Neighbors Accuracy: 0.62
Support Vector Machine Accuracy: 0.73
```




8. Feature Importance

```
# 9.1 Using Random Forest
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
importances = rf.feature_importances_


# Plot
feature_names = df.drop('DiseasePresent', axis=1).columns
feat_df = pd.Series(importances, index=feature_names).sort_values(ascending=False)
feat_df.plot(kind='bar', title="Feature Importance (Random Forest)")
plt.show()
```

```
# 9.2 Drop least important feature (if needed)
# For demo, dropping the least important one
X_reduced = df.drop(columns=['DiseasePresent', feat_df.idxmin()])
X_reduced_scaled = scaler.fit_transform(X_reduced)
```



9. Confusion Matrix

```
y_pred_rf = rf.predict(X_test)
cm = confusion_matrix(y_test, y_pred_rf)
sns.heatmap(cm, annot=True, fmt='d')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



10. Classification Metrics

```
print("11.1 Classification Report:")
print(classification_report(y_test, y_pred_rf))

print("11.2 Accuracy:", accuracy_score(y_test, y_pred_rf))
```

```
print("11.3 Error:", 1 - accuracy_score(y_test, y_pred_rf))
print("11.4 Precision:", precision_score(y_test, y_pred_rf))
print("11.5 Recall:", recall_score(y_test, y_pred_rf))
```

```
# 11.6 TPR (Same as Recall)
tpr = recall_score(y_test, y_pred_rf)
print("11.6 True Positive Rate:", tpr)
```

```
# 11.7 FPR
fp = cm[0][1]
tn = cm[0][0]
fpr = fp / (fp + tn)
print("11.7 False Positive Rate:", fpr)
```

```
# 11.8 Specificity
specificity = tn / (tn + fp)
print("11.8 Specificity (TNR):", specificity)
```

```
# 11.9 F1 Score
print("11.9 F1 Score:", f1_score(y_test, y_pred_rf))
```

```
# 11.10 Support (from classification report)
```

```
→ 11.1 Classification Report:
```

	precision	recall	f1-score	support
0	0.73	0.94	0.82	147
1	0.18	0.04	0.06	53
accuracy			0.70	200
macro avg	0.46	0.49	0.44	200
weighted avg	0.58	0.70	0.62	200

```
11.2 Accuracy: 0.7
11.3 Error: 0.30000000000000004
11.4 Precision: 0.18181818181818182
11.5 Recall: 0.03773584905660377
11.6 True Positive Rate: 0.03773584905660377
11.7 False Positive Rate: 0.061224489795918366
11.8 Specificity (TNR): 0.9387755102040817
11.9 F1 Score: 0.0625
```

11. Cross-Validation

```
cv_scores = cross_val_score(RandomForestClassifier(), X_scaled, y, cv=5)
print("12. Cross-validation scores:", cv_scores)
print("Mean CV Accuracy:", np.mean(cv_scores))
```

```
→ 12. Cross-validation scores: [0.7  0.69 0.69 0.695 0.695]
Mean CV Accuracy: 0.694
```

12. Results and Conclusion

```
print("\n13. Results & Conclusion:")
best_model = max(accuracy_scores, key=accuracy_scores.get)
print(f"Best model: {best_model} with accuracy: {accuracy_scores[best_model]:.2f}")
```

```
→ 13. Results & Conclusion:
Best model: Logistic Regression with accuracy: 0.73
```

13. References

```
print("\n14. References:")
print("- Scikit-learn documentation: https://scikit-learn.org/")
print("- Seaborn and Matplotlib for visualization")
print("- Data simulated for demo purposes")
```